

Statistical implications of clipped Hebbian learning of cell assemblies.¹

Andreas Knoblauch^a

^a*Department of Neural Information Processing
University of Ulm, Oberer Eselsberg, D-89069 Ulm, Germany
email: knoblauch@neuro.informatik.uni-ulm.de*

Abstract

In this study binary associative networks of the Willshaw type are analyzed with respect to the effect of clipped Hebbian learning on the distribution of postsynaptic potentials when stimulating with random activation patterns. It is shown that the variance in the postsynaptic potentials grows with the square of the stimulation strength if the synapses have been generated by Hebbian learning of many overlapping patterns, but only linearly for independent random synapses. This result bears implications both for analysis of associative memory and the detection of Hebbian cell assemblies in neurophysiological experiments.

Key words: cell assemblies, Hebbian learning, associative memory, Willshaw model, postsynaptic potentials

1 Introduction

Although cell assemblies have been postulated by Donald Hebb more than half a century ago [4], so far they have not yet been proven (or disproven) to occur in the real brain. This is mainly because of difficulties in recording simultaneously from a large number of single neurons with high spatial and temporal resolution. This study suggests an alternative test by analyzing statistical properties of the neurons' postsynaptic potentials after random stimulation in a simple binary network model [14,15]. It is shown that the variance in the postsynaptic potentials grows with the square of the stimulation strength if the synapses have been generated by Hebbian learning of many overlapping cell assemblies, but only linearly for independently generated random synapses.

¹ The author has been supported by the MirrorBot project of the European Union.

entries in the address pattern \tilde{u} occur also in one of the original patterns, u^μ , then the one-entries in the retrieval result \hat{u} will always be a superset of the ones in the original pattern. Indeed, this strategy is the only possible choice if one assumes that the address pattern contains no ‘false alarms’, and it plays also an important role for pattern separation in spiking associative memories with time-continuous retrievals (cf. [7,6]).

Generally, the probability of a retrieval error will increase with the fraction p_1 of active synapses. p_1 is also referred to as the *matrix load* and increases with the number M of stored patterns. For random patterns we obtain

$$p_1 \approx 1 - (1 - k^2/n^2)^M \approx 1 - e^{-Mk^2/n^2}, \quad (2)$$

It is the main matter of the theory of neural associative memory to determine how much information and how many patterns can be stored safely in a network of n neurons [9,12,5,6]. Given an address pattern, the general strategy is to divide the neuron population into “correct” and “false” neurons, and then to determine the two distributions of membrane potentials for the two classes. Finally, applying an optimally chosen threshold Θ extracts ideally exactly the “correct” neurons.

A popular approximation is to assume that the activated synapses would have been generated independently of each other which neglects the fact that storing a pattern will activate a block of k^2 synapses at a time. However, applying this *binomial approximation*, analysis becomes relatively easy. For example, let us assume that the address pattern \tilde{u} contains $\lambda \cdot k$ one-entries of a previously stored pattern u^μ , and $\kappa \cdot k$ randomly chosen noisy one-entries ($0 < \lambda \leq 1$, $\kappa \geq 0$), such that the total address pattern activity is $z := (\lambda + \kappa) \cdot k$. Using the binomial probability,

$$p_B(x; N, p) := \binom{N}{x} \cdot p^x \cdot (1 - p)^{N-x}, \quad (3)$$

the membrane potential X_c of a “correct” neuron, and the membrane potential X_f of a “false” neuron are approximately distributed according to the probabilities

$$\text{pr}[X_c = x] \approx p_B(x - \lambda \cdot k; \kappa \cdot k, p_1), \quad (4)$$

$$\text{pr}[X_f = x] \approx p_B(x; z, p_1). \quad (5)$$

For $\kappa = 0$, for example, we can simply apply the Willshaw threshold strategy ($\Theta = z = \lambda \cdot k$) and write for the error probability that a given “false” neuron gets activated, $p_{01} = \text{pr}[X_f = z] \approx p_B(z; z, p_1) = p_1^z$.

Unfortunately, it turns out that the binomial approximation can be quite bad, both qualitatively and quantitatively [6]. Figure 2 illustrates two phenomena that can occur for the potential distributions: (1) Oscillatory modulations and (2) massive underestimation of the variance, in particular for large address pattern activity z .




Fig. 2. The binomial approximation of the neuron potential distribution can be bad. **a:** Simulations reveal oscillatory modulations in the neuron potentials (black) not captured by the binomial approximation (gray). **b:** The s.d. σ_B of the binomial approximation can underestimate massively the true s.d. σ_W of the Willshaw distribution. *Crosses* refer to simulation experiments, *solid line* uses eqs. 14, 15 (cf. [6]).

Both phenomena can be explained by the exact formula $\text{pr}[X_f = x] = p_W(x; n, k, M, z)$ found by Buckingham and Willshaw ([3,2]; see also [13,8] and the appendix in [9]),

$$p_W(x; n, k, M, z) := \sum_{i=0}^M p_B(i; M, \frac{k}{n}) \cdot p_B(x; z, 1 - (1 - \frac{k}{n})^i), \quad (6)$$

which, however, cannot be evaluated and analyzed as easily as the binomial approximation. In the following p_W is referred to as the *Willshaw distribution*. First note that p_W is a superposition of M binomial probabilities $p_B(x; \dots)$. This explains the oscillatory modulations which occur if the standard deviation (s.d.) of the component binomials is small compared to the distances of the means (see [6] for details). In the following section I will determine a good approximation for the variance of p_W which explains the underestimation by the binomial approximation.

3 The variance of the Willshaw distribution

Now we try to approximate the variance $\text{Var}(X_f) = E(X_f^2) - E(X_f)^2$ from the formula of Buckingham and Willshaw (eq. 6). First note that the binomial approximation already gives us the correct expectation (cf. eq. 2)

$$E(X_f) = z \cdot p_1 = z \cdot (1 - (1 - \frac{k^2}{n^2})^M) \quad (7)$$

In order to infer the variance of the Willshaw distribution from eq. 6 the following equations turn out to be useful,

$$\sum_{i=0}^M p_B(i; M, p) \cdot (1-p)^i = (1-p^2)^M \quad (\approx e^{-Mp^2}), \quad (8)$$

$$\sum_{i=0}^M p_B(i; M, p) \cdot q^i \approx e^{-rMp^2(1-p(r-1)/2)}. \quad (9)$$

The second approximation requires $0 < p \ll 1$ and $q = (1-p)^r$, where the sum over p_B has been replaced by an integral over a Gaussian (cf. section 3.6.3 in [6]). With this we can infer from eq. 6 the second moment of the Willshaw distribution

$$E(X_f^2) = \sum_{x=0}^z x^2 \cdot \sum_{i=0}^M p_B(i; M, \frac{k}{n}) \cdot p_B(x; z, 1 - (1 - \frac{k}{n})^i) \quad (10)$$

$$= z^2 - z(2z-1) \cdot (1 - \frac{k^2}{n^2})^M + z(z-1) \cdot \sum_{i=0}^M p_B(i; M, \frac{k}{n}) \cdot (1 - \frac{k}{n})^{2i} \quad (11)$$

$$\approx z^2 - z \cdot (2z-1) \cdot (1 - \frac{k^2}{n^2})^M + z \cdot (z-1) \cdot e^{-2M \frac{k^2}{n^2} (1 - \frac{k}{2n})} \quad (12)$$

Finally, for the variance σ_W^2 of the Willshaw distribution the following approximation can be found,

$$\sigma_W^2 := \text{Var}(X_f^2) = E(X_f^2) - E(X_f)^2 \quad (13)$$

$$\approx z \cdot p_1 \cdot (1-p_1) - (z^2 - z) \cdot \frac{k}{n} \cdot (1-p_1)^2 \cdot \ln(1-p_1) \quad (14)$$

$$\approx z \cdot p_1 \cdot (1-p_1) =: \sigma_B^2. \quad (15)$$

The first approximation is quite good (cf., Fig. 2b). The second approximation yields the variance σ_B^2 suggested by the classical binomial approximation, and is justified only for small k/n and z (cf. [6]). Thus, the classical analysis of binary associative nets [15,9,12,5] may overestimate the storage capacity and fault tolerance, in particular if the address pattern contains a large number z of one-entries (cf. [6]). On the other hand, this effect could be exploited in a neurophysiological experiment to verify the hypothesis of local cell assemblies in a cortical column. The basic idea would be to stimulate randomly a large neuron number z , and then to compute the coefficient of variation $\text{CV} := \text{s.d.}/\text{mean}$ for the amplitude of the postsynaptic potentials (or currents). If $\text{CV} \rightarrow 0$ for increasing z , this would falsify the assembly hypothesis, while an asymptotically positive CV would support it.

4 Conclusion and discussion

In order to explain some discrepancies between the true membrane potential distributions and a simple binomial approximation widely used in the literature [15,9,12,5] (for exact approaches see [3,2,13] and the appendix in [9]), a good approximation for the variance of the Willshaw distribution of membrane potentials has been derived (cf. Fig. 2 and eqs. 14 and 15). It turned out that the variance increase is $\sim z^2$ with the stimulation strength z , while the binomial approximation suggests only a linear growth $\sim z$. This result has two important implications (cf. [6]): (i) It improves the classical analysis of binary associative networks which use the binomial approximation and therefore overestimate fault tolerance and storage capacities. (ii) It suggests a neurophysiological experiment which could contribute to prove or disprove the occurrence of local cell assemblies in the cortical areas of the brain. Of course, the analyzed model is very abstract compared to the neurobiological reality (cf. [6,1]). Whether this approach really leads to viable experiments remains to be seen in future analytical work and discussions with neurophysiologists.

References

- [1] V. Braitenberg and A. Schüz. *Anatomy of the cortex. Statistics and geometry*. Springer-Verlag, Berlin, 1991.
- [2] J.T. Buckingham. Delicate nets, faint recollections: a study of partially connected associative network memories. *PhD thesis, University of Edinburgh*, 1991.
- [3] J.T. Buckingham and D.J. Willshaw. Performance characteristics of associative nets. *Network: Computation in Neural Systems*, 3:407–414, 1992.
- [4] D.O. Hebb. *The organization of behavior. A neuropsychological theory*. Wiley, New York, 1949.
- [5] A. Knoblauch. Optimal matrix compression yields storage capacity 1 for binary Willshaw associative memory. In O. Kaynak, E. Alpaydin, E. Oja, and L. Xu, editors, *Artificial Neural Networks and Neural Information Processing - ICANN/ICONIP 2003*, LNCS 2714, pages 325–332. Springer Verlag, Berlin, 2003.
- [6] A. Knoblauch. Synchronization and pattern separation in spiking associative memory and visual cortical areas. *PhD thesis, Department of Neural Information Processing, University of Ulm, Germany*, 2003.
- [7] A. Knoblauch and G. Palm. Pattern separation and synchronization in spiking associative memories and visual areas. *Neural Networks*, 14:763–780, 2001.

- [8] J.-P. Nadal. Associative memory: on the (puzzling) sparse coding limit. *J.Phys. A: Math. Gen.*, 24:1093–1101, 1991.
- [9] G. Palm. On associative memories. *Biological Cybernetics*, 36:19–31, 1980.
- [10] G. Palm. *Neural Assemblies. An Alternative Approach to Artificial Intelligence*. Springer, Berlin, 1982.
- [11] G. Palm. Cell assemblies as a guideline for brain research. *Concepts in Neuroscience*, 1:133–148, 1990.
- [12] G. Palm. Memory capacities of local rules for synaptic modification. A comparative review. *Concepts in Neuroscience*, 2:97–128, 1991.
- [13] F.T. Sommer and G. Palm. Improved bidirectional retrieval of sparse patterns stored by hebbian learning. *Neural Networks*, 12:281–297, 1999.
- [14] K. Steinbuch. Die Lernmatrix. *Kybernetik*, 1:36–45, 1961.
- [15] D.J. Willshaw, O.P. Buneman, and H.C. Longuet-Higgins. Non-holographic associative memory. *Nature*, 222:960–962, 1969.